

Heuristic Analysis

Objective

The objective of this work is to generate, implement, analyze, and compare three custom heuristic functions used to calculate the value of an *Isolation* board. These functions are compared against the *Improved* heuristic while being executed by an agent using the *Iterative deepening* search algorithm. This agent is called the **ID_Improved** agent.

Both agents are executed in a tournament against other three agents each of them using a different search algorithm and with three different heuristic functions:

- **Random agent:** It is an agent that chooses randomly. This agent is not executed with any heuristic function (as it does not use any).
- **Minimax Agent (MM):** It is an agent that implements a fixed-depth minimax search algorithm.
- **Alpha-Beta agent (AB):** It is an agent that implements the fixed-depth Alpha-Beta pruning search algorithm.

The **AB** and **MM** agents are evaluated with three different heuristics:

- **Null heuristic:** This heuristic presumes no knowledge for non-terminal states, and returns the same uninformative value for all other states.
- **Open move heuristic:** This heuristic implements the basic evaluation function described in lecture that outputs a score equal to the number of available moves for the player.
- **Improved heuristic:** This heuristic calculates the difference of available moves between the two players.

The tournament estimates the strength rating of the student-agent with iterative deepening and a custom heuristic evaluation function against the **AB** and **MM** test agents by running a round-robin tournament.

The *student agent* and the *ID_Improved* agents play 20 "fair" matches against each test agent. The matches are fair because, for each match, the board is initialized randomly for both players, and the players play each match twice -- switching the player order between games. This helps to correct for imbalances in the game due to both starting position and initiative.

The results for the *ID_Improved* agent were pretty stable on the computer used to execute the tournaments. An example of a typical result is the following:

```
[-] Match 1: ID_Improved vs Random      Result: 40 to 0
    Match 2: ID_Improved vs MM_Null      Result: 38 to 2
    Match 3: ID_Improved vs MM_Open      Result: 30 to 10
```

Match 4:	ID_Improved vs MM_Improved	Result: 25 to 15
Match 5:	ID_Improved vs AB_Null	Result: 31 to 9
Match 6:	ID_Improved vs AB_Open	Result: 28 to 12
Match 7:	ID_Improved vs AB_Improved	Result: 26 to 14
ID_Improved		77.86%

This agent played in total 4200 games and won 3257 of them giving it a total score of 77.54%

Heuristics Analyzed

Available moves in common

This heuristic function calculates the size of the intersection of the sets of open moves of each player. The board score is positive if the active player is the current player and negative otherwise.

The rationale behind this heuristic is that the moves in common are the moves were the active agent can close options for the opponent possibly creating a partition. Thus, when the active player is our agent the score is positive meaning that from this board some good moves can be possible done and viceversa.

A tipical tournament for this agent was as follows:

Match 1:	Student vs Random	Result: 30 to 10
Match 2:	Student vs MM_Null	Result: 24 to 16
Match 3:	Student vs MM_Open	Result: 24 to 16
Match 4:	Student vs MM_Improved	Result: 15 to 25
Match 5:	Student vs AB_Null	Result: 25 to 15
Match 6:	Student vs AB_Open	Result: 21 to 19
Match 7:	Student vs AB_Improved	Result: 17 to 23
Student		55.71%

This agent played in total 1400 games and won 777 of them giving it a total score of 55.50% which is extremelly lower than the *ID_Improved* agent's score.

Distance to the opponent

This heuristic function calculates the distance between the two players.

The rationale behind this heuristic is that it could be better to be further away from the opponent and prevent the blocking of movements.

A tipical tournament for this agent was as follows:

Match 1:	Student vs Random	Result: 32 to 8
Match 2:	Student vs MM_Null	Result: 26 to 14
Match 3:	Student vs MM_Open	Result: 17 to 23
Match 4:	Student vs MM_Improved	Result: 19 to 21
Match 5:	Student vs AB_Null	Result: 24 to 16

```

Match 6:   Student   vs   AB_Open   Result: 22 to 18
Match 7:   Student   vs   AB_Improved Result: 18 to 22
Student                                     56.43%

```

This agent played in total 1400 games and won 789 of them giving it a total score of 56.35% which is better than the *common moves* agent, but still a lot worse than the *ID_Improved* agent.

Linear combination of moves

The heuristics presented in the lectures are based on the number of available moves that each agent has with a preference for the boards that have more available movements for the player. It is suggested that other combinations can be implemented assigning different weights to each player moves creating a family of heuristics generated by the linear combination of the variables:

```

x - amount of open moves of the player
y - amount of open moves of the opponent

```

Each heuristic has the form:

```

h = a*x + b*y + c

```

It is possible to perform a grid search over the values of a,b,c to find the values that maximize the heuristic score.

A grid search was performed restricting the possible values of a,b, and c to following ranges:

```

-3 <= x,y <= 3
c in {-24, -16, -8, 0, 8, 16, 24}

```

The selected ranges are mostly arbitrary and were kept small in order to be able to execute the search in reasonable time.

The evaluation of each heuristic was performed by running a tournament of 100 matches against the *AB* agent and counting the winning matches. The results of the whole grid search can be found in the file `grid_search_results.txt` (https://github.com/edhzsz/AIND-Isolation/blob/master/grid_search_results.txt) in this repository.

The best values found are: a = 1 b = -2 c = -8 Which got a score of 75.33%.

It is worth noting that the *open moves* and *improved* heuristics described in the lectures were also evaluated and they got a score of 67.66% and 68.33% respectively.

An agent was generated and tested using this values and the results of a typical tournament for this agent are as follows:

```

Match 1:   Student   vs   Random   Result: 37 to 3
Match 2:   Student   vs   MM_Null   Result: 37 to 3
Match 3:   Student   vs   MM_Open   Result: 31 to 9
Match 4:   Student   vs   MM_Improved Result: 29 to 11

```

Match 5:	Student	vs	AB_Null	Result: 35 to 5
Match 6:	Student	vs	AB_Open	Result: 27 to 13
Match 7:	Student	vs	AB_Improved	Result: 29 to 11
			Student	80.36%

This agent played in total 1400 games and won 1119 of them giving it a total score of 79.92% which is better than the *ID_Improved* agent's score making it suitable to be used as the evaluation function by an agent that plays the Isolation game with horses.